

---

## Reconfigurable Hardware Technology

An FPGA is an integrated circuit that belongs to a family of programmable devices called Programmable Logic Devices (PLDs). An FPGA contains tenths of thousands of building blocks, known as *Configuration Logic Blocks* (CLB) connected through programmable interconnections. Those CLBs can be reconfigured by the designers themselves resulting in a functionally new digital circuit, this way, virtually any kind of digital circuit can be implemented using FPGAs [11, 272, 304, 244].

At first, FPGA devices were mainly applied for logic design, and as a consequence of that, numerous tools were designed for synthesizing logic designs on them. Among those tools, Hardware Description Languages (HDL) and schematic diagram editors have been traditionally used as a starting point for such a synthesis process. Among the many hardware description languages available today, Verilog, and especially, VHDL, are the two most widely spread hardware languages.

In recent years, FPGAs have been used for reconfigurable computing when the main goal is to obtain high performance at a reasonable cost out of hardware implemented algorithms. The main advantage of FPGAs is their reconfigurability, i.e., they can be used for different purposes at different stages of a computation and they can be, at least partially, reprogrammed on run-time. The two most popular FPGA manufacturers are Xilinx [396] and Altera [4]. Those two makers have over 70% of the FPGA market share.

Besides Cryptography, applications of FPGAs can be found in the domains of evolvable and biologically-inspired hardware, network processors, real-time systems, rapid ASIC prototyping, digital signal processing, interactive multimedia, machine vision, computer graphics, robotics, embedded applications, and so forth. In general, FPGAs tend to be an excellent choice when dealing with algorithms that can benefit from the high parallelism offered by the FPGA fine-grained architecture.

In this chapter we present the generalities of FPGA technology. We stress that the material of this Chapter is mainly intended for those readers non-familiar with this technology.

We begin in Section 3.1 by reviewing some historical milestones of FPGA development and then we review in Section 3.2 the two most currently used FPGA technologies, namely, Xilinx and Altera. Then we compare in Section 3.3 the performance of FPGA realizations against the ones on ASICs and general-purpose processor platforms. We continue in Section 3.4 by briefly introducing the reconfigurable computing paradigm main concepts. In Section 3.5 we review several key strategies to achieve good designs for cryptographic applications. Then, we define in Section 3.6 several metrics and figures of merit needed to evaluate design performance for reconfigurable computing as well as several security concerns related to FPGA technology. In Section 3.7 we give a brief overview of some of the security concerns and attacks on FPGA technology. Finally, in Section 3.8 concluding remarks are given.

More experimented readers might be interested in reviewing more advanced material. For them, we recommend excellent sources such as the ones found in [124, 365, 217, 199, 192]. Those readers having more technology oriented interests may profit from consulting [259, 244] as well.

### 3.1 Antecedents

The concept of reconfigurable computing was first introduced by G. Estrin in 1960 [101]. His invention consisted of a hybrid machine composed by a general purpose microprocessor interconnected with programmable logic devices. The programmable logic could be configured for accomplishing a specific function with the characteristic efficiency of hardware designs. Once the function was completed, another task could be performed by manually reconfiguring the hardware. This resulted in a hybrid computer structure combining the best features of software (flexibility) and hardware (speed) platforms. It is nothing but remarkable how Estrin's concept come close to what is offered by nowadays modern reconfigurable devices [217].

In the mid 1970s, Programmable Logic Devices (PLDs) were introduced by companies such as IBM, Monolithic Memories, Inc (MMI) and AMD. The first PLDs were called PAL (Programmable Array Logic) or PLA (Programmable Logic Array) depending on the programming scheme utilized [272]. Earlier PLDs consisted of logic gate arrays with no clocked memory components. However, *registered* PLDs including one flip-flop at each output of the circuit, were soon available. Register PLDs allowed for the first time the design of simple reprogrammable sequential circuits.

An innovation of PAL devices was the Generic Array Logic (GAL) device, which had the same logical properties as the PAL but the functionality could be erased and reprogrammed. From the point of view of today's standards,

PALs and GALs devices are small devices having an equivalent computational power of just some few hundred logic gates.

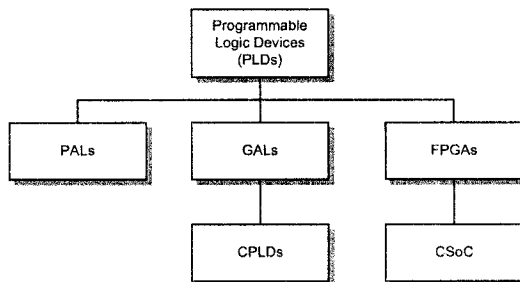
As a consequence of Moore's law, the semiconductor technology has experienced an unrelenting improvement over the last three decades. That allowed the integration in the mid 1980s of several either GAL or PAL devices on the same chip, thus given birth to the CPLD (Complex PLD) devices. CPLDs can emulate the computational power of hundreds of thousands of logic gates and they are still very popular due to their outstanding cost-benefit compromise (some CPLD devices can be bought for less than a dollar). A typical modern CPLD device has a structure consisting of several GAL blocks whose outputs are connected to a switch matrix used for programming the interconnections as well as the Input/Output pins. Each GAL block consists of one or more programmable sum-of-products logic arrays ended with a relatively small number of registers. CPLDs are usually programmed via a serial data port that can be connected to a personal computer. Internally, the CPLD contains a decoding module that interprets the data stream in order to perform a specific logic function. The preferred standard for this programming method is the IEEE 1149.1 standard usually known as Joint Test Action Group (JTAG) interface [272]. As of 2006, most CPLDs are non-volatile electrically-erasable programmable devices.

Field Programmable Gate Array (FPGA) devices were introduced by Xilinx in the mid 1980s. Roughly speaking, FPGA devices are built using a grid of logic gates. They differ from CPLDs in several key aspects. FPGA architectures consists of a matrix of Configurable Logic Blocks (CLBs) interconnected by an intricate array of switch matrices. This architecture provides great flexibility to hardware designers but it also implies much more sophisticated routing technologies [123]. The fact that most modern FPGAs have higher-level embedded modules such as built-in multipliers, distributed RAM blocks and so on is another important difference with CPLD devices. Moreover, in contrast to CPLD devices, most modern FPGAs support (at least partially) *in-system* reconfiguration, thus allowing designs to be changed dynamically "on run-time". This feature can be particularly useful for system updates.

Significant technical advances have led to architectures that combine FPGA's logic blocks and interconnect matrices, with one or more microprocessors and memory blocks integrated on a single chip. This hybrid technology is called Configurable System-on-Chip (CSoC). Examples of the CSoC technology are the Xilinx Virtex-II PRO, and the Virtex-4 and Virtex-5 FPGA families, which include one or more hard-core PowerPC processors embedded along with the FPGA's logic fabric [398, 396, 397].

Alternatively, *soft* processor cores that are implemented using part of the FPGA logic fabric are also available. This approach is more flexible and less costly than the CSoC technology [217]. Many soft processor cores are now available in commercial products. Some of the most notorious examples are: Xilinx 32-bit MicroBlaze and PicoBlaze, and the Altera Nios and the 32-bit

Nios II processor [394, 5]. These soft processor cores are configurable in the sense that the designer can introduce new custom instructions or processor data paths. Furthermore, unlike the hard-core processors included in the CSoC technology, designers can add as many soft processor cores as they may need (some designs could include 64 such processors or even more [130, 217]).



**Fig. 3.1.** A Taxonomy of Programmable Logic Devices

Fig. 3.1 shows the taxonomy of the programmable logic devices just discussed. In the next Section, more specific details of the FPGA device internal architecture are given.

### 3.2 Field Programmable Gate Arrays

In a very rough way, an FPGA can be seen as a matrix of Configurable Logic Blocks (CLBs), where not only the logic but also the connection is user programmable. The specific design of the CLB blocks varies from manufacturer to manufacturer and even, from device to device. A CLB can be as simple as just one four-input Look Up table (LUT) or as complex as a 4-input Arithmetic Logic Unit (ALU), or a 6-input Look Up Table [398]. It is customary to define the *granularity* of the reconfigurable logic as the size of the smallest functional unit that can be addressed by the programming tools.

Architectures having finer granularity tend to be more useful for data manipulation at bit level and, in general, for combinatorial circuits. On the other hand, blocks with a *coarse grain* granularity are better suited for higher levels of data manipulation, for example, for developing circuits at register-transfer level. The level of granularity has a great impact in the device configuration time. Indeed, a device with low granularity (also known as fine-grained devices) requires many configuration points producing a bigger vector data for reconfiguration. That extra routing has an unavoidable cost on power and area.

On the other hand, a coarse grained architecture tends to decrease its performance when dealing with computations smaller than what its granularity

is. For example, if for a specific application, bit-level operations are required and the smallest functional unit is four-bit wide, then a waste of three bits would occur.

FPGA interconnection has a major role in the performance of an FPGA device due to the need of fast and efficient communication highways among the different logic blocks which are organized by rows and columns. Xilinx devices<sup>1</sup> are equipped with four kinds of interconnects: long lines, hex lines, double lines and direct lines. Direct connect lines are intended for connecting neighbor components (for example, carry circuitry). Hex and double lines are medium length interconnects aimed for connecting many CLBs. Finally long lines interconnects are implemented along the whole chip and are normally utilized for global system signals.

In recent years, huge technological developments have had a great impact on FPGA industry. The most advanced FPGA devices operate up to 550 MHz internal clock with a gate complexity of over 10 Million gates on a single Virtex-5 FPGA chip using a technology of just 65 nm operating at 1.0V [395]. The improvements in technology are not only limited to an ever growing internal number of logic gates but also to the addition of many functional blocks like fast access memories, multipliers or even microprocessors integrated within the same chip.

There are quite a few FPGA commercial manufacturers, and usually each one of them has developed one or more device families. Table 3.1 shows some of the most popular manufacturer families.


**Table 3.1.** FPGA Manufacturers and Their Devices

Manufacturer	FPGA Family	Feature
Xilinx	Virtex-5, Virtex-4, VirtexII, Spartan III	FPGA market leader 65nm technology
Altera	Stratix, Stratix II, Cyclone	90nm technology
Lattice	LatticeXP	first non-volatile FPGA
Actel	Fusion, M7Fusion	first mixed-signal FPGA
Quick Logic	Eclipse II	programmable-only-once FPGA
Atmel	AT40KAL	fine-grained reconfigurable
Achronix	Achronix-ULTRA	1.6GHz - 2.2GHz speed

### 3.2.1 Case of Study I: Xilinx FPGAs

Table 3.2 shows the main features that are included in the Xilinx FPGA families: Virtex-5, Virtex-4, Virtex II Pro and Spartan 3E. The architecture of those Xilinx FPGA families consists of five fundamental functional elements,

<sup>1</sup> At the time that this book was being written, Xilinx released the Virtex-5 family which has a radically different CLB interconnection pattern [395].



**Fig. 3.2.** Xilinx Virtex II Architecture

**Table 3.2.** Xilinx FPGA Families Virtex-5, Virtex-4, Virtex II Pro and Spartan 3E

Feature/family	Virtex-5	Virtex-4	Virtex II Pro	Spartan 3 & 3E
Logic Cells	up to 330K	12K-200K	3K-99K	1.7K-74K
BRAM (18Kbits each)	576	36-512	12-444	4-104
Multipliers	32 – 192 <sup>1</sup>	32-512	12-444	4-104
DCM	up to 18	4-20	4-12	2-18
IOBs	up to 1200	240-960	204-1164	63-633
DSP Slices	32-192	32-192	—	—
PowerPC Blocks	N/A	0-2	0-2	—
Max. freq.	550MHz	500MHz	547 MHz	up to 300MHz
Technology	1.0V, 65nm copper CMOS	1.2V, 90nm, triple-oxide process	1.5V, 130nm, 9-layer CMOS	1.2V, 90nm, triple-oxide process
Price	N/A	From \$345	From \$139	From \$2 up to \$85

<sup>1</sup>25 × 18 embedded multipliers

- Configurable Logic Block (CLB) and Slice architecture;
- Input/Output Blocks (IOBs);
- Block RAM;
- Dedicated Multipliers and;
- Digital Clock Managers (DCMs).

Those components are physically organized in a regular array as shown in Fig. 3.2. In the following we explain each one of those five elements<sup>2</sup>.

<sup>2</sup> Virtex-5 devices can be considered second generation FPGA devices. In particular, a Virtex-5 slice contains four true 6-input Look Up Tables (LUTs).




Fig. 3.3. Xilinx CLB


### Configuration Logic Blocks (CLBs)

The Configurable Logic Blocs (CLBs) are the most important and abundant hardware resource of an FPGA. They are typically utilized for both, combinatorial and synchronous logic design. Each CLB is composed of four *slices*<sup>3</sup>, which are interconnected as shown in Fig. 3.3. The slices are grouped by pairs and each pair is organized by a column with independent carry chain [395].


All four slices have the following common elements: two Look-Up Tables (LUTs), two type D flip-flops, multiplexers, logic circuits for carry handling and arithmetic logic gates. Both, the left and right pair of slices utilize those elements for providing logic functions, arithmetic and ROM. Besides that, the left pair supports two additional functions: data storage using a distributed RAM and 16-bit shift register functionality. Fig.3.4 shows the internal structure of a CLB. The atomic building block of a Virtex CLB is the logic cell (LC). An LC includes the Look-Up Table block, carry logic, and a storage element (flip-flop) as shown in Figure 3.5.

As it was mentioned, a CLB can be configured to work into two modes: logic mode and memory mode. As shown in Fig. 3.6, in logic mode, each CLB Look Up Table behaves as a combinational logic block and a one bit register. In the case of Xilinx devices those Look Up Tables can be reprogrammed to any arbitrary combinational logic function of four inputs/one output. In memory mode, Look Up Table blocks behave as two small pieces of memory blocks.


<sup>3</sup> Slice is a term introduced by Xilinx. It specifies a basic processing unit in a Xilinx FPGA.



**Fig. 3.4.** Slice Structure



**Fig. 3.5.** VirtexE Logic Cell (LC)



**Fig. 3.6.** CLB Configuration Modes

**Input/Output Blocks**

Input/output Blocks (IOB) provide a bidirectional programmable interface between the outside world and the internal logic structure of the FPGA device.



There exist three types of routing possibilities for an IOB: output signal, input signal and third state (high impedance) signal. Each one of those signals has their own pair of storage elements that can behave as registers or as latches [395].

**Block RAM**

Virtex devices include built-in 18K-bit RAM memory, called BRAM. BRAMs can be configured in a synchronous manner. BRAMs are intended for storing big amounts of data, while the distributed RAM is more useful for storing small amounts of data.

BRAMs are polymorphic blocks in the sense that its width and depth can be configured. Even multiple blocks can be connected in a back-to-back configuration in order to create wider and/or deeper memory blocks. A BRAM block supports several configuration modes, including single or double port RAM and several possible combination of data/address sizes as is shown in Table 3.3.

**Table 3.3.** Dual-Port BRAM Configurations

Configuration	Depth	Data bits	Parity bits
16K × 1 bit	16Kb	1	0
8K × 2 bit	8Kb	2	0
4K × 4 bit	4Kb	4	0
2K × 9 bit	2Kb	8	1
1K × 18 bit	1Kb	16	2
512 × 36 bit	512	32	4

**18x18 Bit Multiplier**

Xilinx FPGAs have several dedicated multiplier blocks. Those multipliers accept two 18-bit operands in two's complement form computing their product also in two's complement form. Such multipliers blocks have been optimized for performing at a high speed while their power consumption is kept low when compared with multipliers directly implemented using the CLB resources. The total number of multipliers varies from device to device as is shown in Table 3.2.

**Digital Clock Managers**

Digital Clock Managers (DCMs) provide a flexible control over clock frequency, phase shift and skew. The three most important functions of DCMs are: To mitigate clock skew due to different arrival times of the clock signal,

to generate an ample range of clock frequencies derived from the master clock signal and, to shift the signal of all its output clock signals with respect to the input clock signal.

### 3.2.2 Case of Study II: Altera FPGAs

Altera offers a wide variety of programmable hardware devices which are grouped into four categories [4].

- Complex Programmable Logic Devices(CPLDs)
- Low-Cost FPGAs
- High-density FPGAs
- Structured ASICs

#### CPLDs

Altera's CPLDs include MAX (EPM3032A, EPM3512A) and MAX-II (EPM 240/G, EPM 2210/G) family of devices. They are low complexity, low density and easy to use CPLD family for which software tools can be downloaded from Internet and they are free of cost.

#### Low-Cost FPGAs

Cyclone (EP1C3,EP1C20) and Cyclone-II (EP2C5, EP2C7) family of devices are considered low cost FPGAs. Their main features include embedded DSP blocks, on chip memory modules and support for embedded processor (NIOS).

#### High-Density FPGAs

The category of high density FPGAs from Altera comprises Stratix-II (EP2S15, EP2S180), Stratix (EP1S10, EP1S80), Stratix<sub>GX</sub>-II (EP2SGX30C/D, EP2SG-X130G) and Stratix<sub>GX</sub> (EP1SGX10C, EP1SGX40G) family of devices. Stratix and Stratix-II families are general purpose FPGAs with fast performance, large on-chip memory modules, and DSP blocks. Stratix<sub>GX</sub> and Stratix<sub>GX</sub>-II families, in addition, include integrated transceivers.

#### Structured ASICs


Structured ASICs comprise Hardcopy (HC1S25, HC240) and Hardcopy-II (HC210W, HC240) solutions. They have similar design flow as that of Stratix and Stratix-II respectively. They are low cost structured ASIC solutions with sufficient number of gates supported by all major EDA vendors.

To provide an idea of what kinds of resources are present in Altera FPGA devices, let us discuss the structure of the Stratix family of devices. Detailed

data sheets of Stratix as well as all other Altera devices can be consulted in [4, 207, 208]. The quantitative information presented in this subsection has been extracted from [4]. Table 3.4 provides a quantitative measure of Stratix major resources, while Fig. 3.7 shows the physical distribution of those resources.

**Table 3.4.** Altera Stratix Devices

Feature	Device						
	EP1S10	EP1S20	EP1S25	EP1S30	EP1S40	EP1S60	EP1S80
Logic Elements	10,570	18,460	25,660	32,470	41,250	57,120	79,040
M512 RAM Blocks	94	194	224	295	384	574	767
M4K RAM Blocks	60	82	138	171	183	292	364
M-RAM Blocks	1	2	2	4	4	6	9
Total RAM bits	0.9205M	1.669M	1.945M	3.317M	3.423M	5.215M	7.427M
DSP Blocks	6	10	10	12	14	18	22
Embedded Multipliers	48	80	80	96	112	144	176
PLLs	6	6	6	10	12	12	12
Maximum I/O Pins	426	586	706	726	822	1022	1203



**Fig. 3.7.** Stratix Block Diagram

As shown in Fig. 3.7, the main building blocks in Stratix devices are the following:

- Logic Array Blocks (LABs)

- Memory Blocks
- Digital Signal Processing (DSP) Blocks
- Input/Output Elements (IOEs)
- Interconnects

### Logic Array Blocks (LABs)

LABs are arranged in rows and columns across the device. Each LAB consists of 10 Logic Elements (LE). An LE is the smallest unit in Stratix architecture. It contains four input LUT, carry chain with carry select capability and a programmable register as shown in Fig. 3.8. The LUT serves as a function generator which can be programmed to any function with four variables. By using LAB-wide control signal, a dynamic addition or subtraction mode can also be selected. It is to be noted that number of resources are not fixed for an LAB in all kind of Altera devices. As an example, a LAB in Stratix-II architecture comprises 8 Adoptive Logic Modules (ALM) where each ALM contains a variety of LUT-based resources.

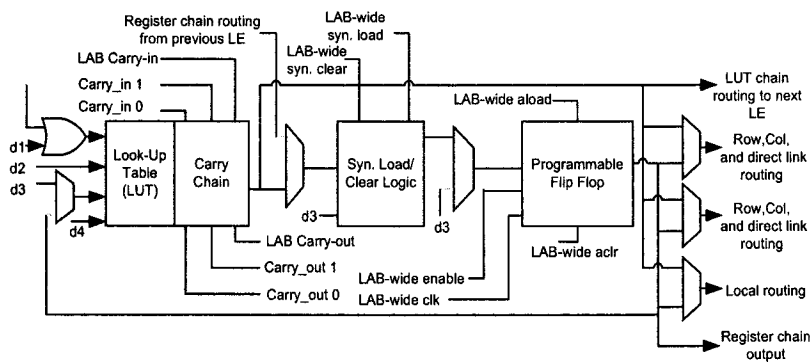


Fig. 3.8. Stratix LE

The Stratix LE can be configured into two modes:

- Normal mode
- Dynamic arithmetic mode

In normal mode, a four input LUT can be used to implement any function. The normal mode is therefore useful for implementing combinational logic and general logic functions. In dynamic arithmetic mode, an LE utilizes four 2-input LUTs which can be mapped to a dynamic adder/subtractor. First two LUTs perform two summations with possible carry-in and the other two LUTs compute carry outputs to drive two chains of the carry select circuitry. The

arithmetic mode is therefore useful for wide range of applications like adders, accumulators, wide parity functions, etc.

### Memory Blocks

Three types of memory blocks are present in Stratix devices as shown in Fig. 3.7. Those are referred to as M512 RAM, M4K RAM and M-RAM (MegaRAM) blocks. M512 RAM is a simple dual port memory with sizes of 512 bits plus parity (576 bits). It can be configured as a maximum 18-bit wide single or dual port memory at up to 318 MHz. M4K is a true dual port memory with 4K bits plus parity. It can be configured as a maximum 36-bit wide dedicated dual port, simple dual or single port memory at 291 MHz. Several M-RAM blocks can also be located individually in logic arrays across the device. It is a true dual port memory with 512K bits plus parity (589,824 bits). A single M-RAM can be configured as a maximum 144-bit wide dedicated dual port, simple dual or single port memory which can operate at 269 MHz.

### DSP Blocks

Those are dedicated Stratix resources which are vertically arranged into two columns in each device. DSP blocks can be configured into either eight  $9 \times 9$ -bit multiplier, four  $18 \times 18$ -bit multiplier or one full  $36 \times 36$  multiplier. In addition, DSP blocks also contain  $18 \times 18$ -bit shift registers, Finite Impulse Response (FIR) and Infinite Impulse Response (IIR) filters.

### Input/Output Elements (IOEs)

Large number of IOEs can be located at the end of LAB row or column around the periphery of a Stratix device as shown in Fig. 3.7. Each I/O element comprises a bi-directional I/O buffer and six registers for buffering input, output and output-enable signals. Each Stratix I/O pin is fed by an I/O element and support several single-ended and differential I/O standards.

### Interconnects

All LEs within the same LAB, or all LABs within the same device or Memory blocks or DSP blocks can be interconnected. A single LE can drive 30 other LEs through locally available fast and direct link interconnects. A direct link is also used by adjacent LABs, memory and DSP block to drive LABs local interconnects. The availability of direct links helps in reducing row and column interconnects resulting on higher performance and flexibility.

**Table 3.5.** Comparing Cryptographic Algorithm Realizations on different Platforms

Algorithm	FPGA		ASIC		$\mu$ Processor	
	Throughput	year	Throughput	year	Throughput	year
MD5	5.86 Gbps [156]	2005	2.09 Gbps [312]	2005	1.27Gbps (est)* [31]	1996
SHA-1	0.9 Gbps [67]	2002	2.006 Gbps [312]	2005	0.678Gbps (est)* [31]	1996
DES	21.3 Gbps [301]	2003	10Gbps [381]	1999	0.127Gbps [22]	1997
AES	25.1Gbps [113]	2005	7.5Gbps [303]	2001	0.8Gbps [109]	2004
1024-bit RSA	6.1 mS [6]	2005	1.47mS [210]	2005	22.1mS [294]	2004
ECC (binary)	17.64 $\mu$ S [54]	2006	190 $\mu$ S [313]	2003	475 $\mu$ S [133]	2004
ECC (prime)	3600 $\mu$ S [262]	2001	190 $\mu$ S [313]	2003	325 $\mu$ S [133]	2004

\*Estimated for a 2GHz Pentium IV from the clock cycle count given in [31]

### 3.3 FPGA Platforms versus ASIC and General-Purpose Processor Platforms

Table 3.5 presents a quick performance comparison of several relevant cryptographic algorithms implemented in three different platforms: Reconfigurable hardware devices, ASIC and general purpose processors. We included implementations for hash functions (MD5 and SHA-1), block ciphers (DES and AES) and public key cryptography (RSA and ECC). All those algorithms will be studied in the next Chapters.

Referring to Table 3.5, it is noticed that software implementations are always slower than either, ASIC or FPGA implementations. The performance gap of software implementations is more noticeable for block ciphers and for the binary elliptic curve cryptosystem. On the contrary, the best reported prime elliptic curve cryptosystem is faster than the fastest FPGA design reported in [262].

We stress that the information included in Table 3.5 is intended for a first order comparison. As it has been already mentioned, it is extremely difficult to make fair performance comparisons among designs implemented in different platforms using the different technologies available at the time of their publications. In the rest of this Section we give some more insights about the advantages/disadvantages of implementing a design on reconfigurable hardware compared with other platform options.

#### 3.3.1 FPGAs versus ASICs

Traditionally, in the design of embedded systems, the Application-Specific Integrated Circuit (ASIC) technology has played a major role for providing high performance and/or low cost building blocks necessary for the vast majority of systems during the (usually) large and sinuous design cycle. In 1980 the usage of reprogrammable components was introduced, and short after that the first FPGA device was developed by Xilinx. FPGA devices offer shorter

design cycle because of its ability of providing fast and accurate functionality testing.

However, the relatively high size and power consumption shown by FPGA devices has been the most important drawback of that technology towards an eventual substitution of the virtually ubiquitous ASIC technology. Therefore, historically FPGAs have been utilized primarily for prototyping development.

In recent years, however, FPGA manufacturers have significantly reduced the gap that still exist between FPGA and ASIC technology, paving the way for the utilization of FPGA not only as prototype tools but also as key components of embedded systems or even, becoming the system itself [364, 149, 331, 199].

However, the exact size of the performance gap between FPGAs and ASICs is currently subject of intense analysis and debate. Recently, several experimental results reported in [192], seems to suggest that for circuits designed utilizing the FPGA fabric only (i.e., LUTs and flip flops), an FPGA design is on average 40 times larger, consumes 12 times more dynamic power and it is 3.2 times slower than a standard ASIC implementation. On the other hand, in [364] it was developed a low-power FPGA core which was specially tailored for battery-powered applications such as those found in the automotive industry. The experimental results show that this solution is competitive with similar ASIC solutions.

Undoubtedly, new technological challenges must be faced for both, FPGA and ASIC platforms when the 45 nm and 32 nm technologies come to place. Under this scenario, it is not certain how FPGA new architectures will deal with the power consumption issue. It might be the case that manufacturers would need to trade device performance for a more flexible/predictable device power-consumption [141].

### 3.3.2 FPGAs versus General-Purpose Processors

The speedup that one can expect by implementing an algorithm on an FPGA device rather than using a general purpose processor (i.e. the traditional CPU) has been well documented in the literature [365, 124]. In [124], speedups of one to two orders of magnitude were measured when executing benchmarks applications in the domains of video and image processing. Roughly speaking, the same range of speedups has been confirmed in cryptographic algorithms.

From the qualitative point of view, it is interesting to study the main factors that produce this phenomenon. On the one hand, the typical maximum clock frequency achieved by FPGA designs fall in the range of 20MHz to 100MHz, while embedded microprocessors have frequencies ranging from 300 to 600 MHz and high-end workstation-class processors have frequencies of up to 3.2GHz. Hence, the clock frequency of general-purpose processors is 10-100 times faster than the typical clock frequency found in FPGA designs. On the other hand, there are two factors that help to compensate and even overcome that component, namely,

1. *FPGA Iteration-level parallelism*, obtained by, among others, loop-unrolling, pipeline and sub-pipeline techniques, and;
2. *FPGA Instruction efficiency*, obtained by carefully designed datapaths, the insertion of distributed memory blocks as needed and, taking advantage of the FPGA low granularity, the elimination of several instructions.

Those two factors combine together for obtaining a notable reduction in the total number of clock cycles required by an FPGA implementation. That reduction implies that CPU implementations may require up to 2500 times more clock cycles than that of FPGA implementations [124]. In other words, even though CPU platforms enjoy a much higher operating clock frequency, this factor is not enough for compensating the enormous clock cycle reduction that can potentially be obtained in FPGA platforms.

In the context of Moore's Law, an examination of peak floating-point performance trends for FPGA and CPU platforms is presented in [365]. The author concludes that although CPUs' performance obeys Moore's law (i.e., it doubles every 18 months), FPGA performance is growing at a rate of four times every two years. For applications using the FPGA new functionality (embedded multipliers, RAM blocks, etc.) the performance increase rate may be as high as five times every two years.

### 3.4 Reconfigurable Computing Paradigm

Reconfigurable computing may be defined as computer processing with highly flexible computing fabric. The main idea of reconfigurable computing is to take advantage of the best of two scenarios: flexibility from general purpose computing and speed from reconfigurable logic.

Some of the reconfigurable computing distinguished features when compared to general purpose microprocessors are [123]:

- Due to the inherent fine-grained granularity the parallelism tends to be very high.
- Registers, latches and even distributed RAM blocks can be created and distributed wherever needed by the data path. This characteristic has a tremendous impact on the device performance because reduces unnecessary re-computations and/or memory accesses.
- The amorphous nature (lack of a fixed architecture) of reconfigurable computing devices, allows the designers to tailor design's data path and control flow arbitrarily.

FPGAs can be properly used for rapid prototyping algorithms at hardware level. Considering the restrictions of FPGA devices, desirable FPGA applications should belong to one or more of the categories listed below.

1. Applications that employ only integer arithmetic or at most low precision fixed point arithmetic.



2. Applications that rely on logical operations to make decisions. Comparators, selectors and multiplexers are good examples of that.
3. Applications amenable for being decomposed in independent and pipelined stages.
4. Applications that show regularity in the way they apply a processing.
5. Applications with locality in the interconnection network they require. That means that the application modules should only have interconnections with their neighbors.

Considering FPGA capabilities and limitations some potential applications for FPGAs are:

1. Image processing algorithms such as point type operations (grey scale transformation, histogram equalization, requantization, etc.) and filtering (template matching, window techniques, convolution/correlation, median filtering, etc.) seem to be good candidates for FPGA implementation.
2. Dynamic programming algorithms requiring only integer arithmetic. Dynamic programming is in essence a bottom up procedure in which solutions to all subproblems are first calculated and then these results are used to solve the whole problem. A good example of this approach is the Floyd's shortest path algorithm.
3. Relaxation techniques requiring fixed point arithmetic. The relaxation technique is an iterative approach useful to many problems, which updates in parallel at each point and in each iteration based on the data available in the most recent updating or in the immediate preceding iteration.
4. Associative retrieval operations. Filling and retrieving data by association appears to be a powerful solution to many high volume information processing elements. An associative processing system is very adequate at recognition and recall from partial information and has remarkable error correcting capabilities. The major advantage of associative memory over RAM is its capability of performing parallel search and parallel comparison operations. There are many examples of that kind of applications: pattern matching, artificial intelligence, computer vision, data encoding, compression, and every application maintaining a dictionary data structure.
5. Highly regular and iterative applications with non-standard word lengths. Cryptography is a meaningful example of this kind of applications since it applies basic transformations mostly based on bit-level operations. Those basic operations are performed in long wordlengths starting from 128 bits to up 4096 bits or even in wordlengths non-standard, such as 163 and 233 bits (in the case of public-key cryptography). The basic transformations are repeated iteratively a number of times to process information in stages. In the following chapters we will explain how to take advantage of cryptographic algorithm features for reconfigurable computing.

### 3.4.1 FPGA Programming

The design cycle for programming FPGAs starts with a behavioral description of the design, using either hardware description languages (HDLs) such as VHDL or Verilog or a schematic design entry. Thereafter, the HDL code is compiled in order to produce a *netlist* which represents the mapping of the HDL code to the actual target device hardware resources. After the first compiling step, the netlist is reprocessed in order to perform the place-and-route process whose main goal is to establish how the different design's modules are going to be physically allocated and connected. This will create a binary file which is used for programming or reprogramming the FPGA device. Most designs included in this book have been compiled using the Xilinx Integrated Software Environment (ISE) version 8.1i software [393].

Hardware Description Languages (HDLs) are analogous to other high level languages (C, C++, etc.) with some significant differences. Both types are processed by a compiler, and both of them are function-oriented languages. However they differ in the way that the compiled code is executed. HDL languages are used for formal description of electronic circuits. They describe circuit's operation, its design, and tests to verify its operation by means of simulation. Typical HDL compilers tools [393], verify, compile and synthesize an HDL code, providing a list of electronic components that represent the circuit and also giving details of how they are connected.

### 3.4.2 VHSIC Hardware Description Language (VHDL)

The Very-High-Speed Integrated Circuit Hardware Description Language (VHDL) was created by the US Department of Defense in the early 1980s. In December of 1987, VHDL was adopted as an IEEE Standard [272]. VHDL is a functional language that borrows much of its structure from the programming language Ada along with a set of constructs for supporting the inherent parallelism of hardware designs.

The original version of VHDL, included a wide range of data types such as, logical (bit and boolean), numerical, character and time, plus bit and character. In later versions, the `std_logic` data type was introduced, along with signed and unsigned types to facilitate arithmetical operations, analog and mixed-signal circuit design extensions [367].

Furthermore, the designer can know how his/her HDL instruction was mapped to FPGA components (such as slices, flip-flops, tri-state buffers, etc.). For example, an *if statement* in HDL describes a multiplexer or a flip-flop. It can occur that the frequent use of this statement would insert large number of multiplexers or flip-flops in a circuit, which is functionally correct but may or may not be efficient. As a matter of fact, HDL languages have been designed favoring a hardware designer perspective, in the sense that first the specific hardware architecture should be envisioned, and then an HDL piece of code representing it should be written. If for instance a programmer requires a

flip-flop functionality then he/she should select a suitable flip flop for the design and then he/she can write a code for it. That would generate a list of components for an electronic circuit prior to its implementation providing a designer complete control over available/used FPGA resources.

### 3.4.3 Other Programming Models for FPGAs

Several voices, both from the Academia and Industry sectors, have stated that the main obstacle towards a massive use of reconfigurable computing lies in the difficulty of programming FPGA devices. After all, HDLs were designed primarily from the perspective of designers trying to describe hardware structures, which quite often implies that an FPGA programmer should be primarily a hardware designer.

Considering that, it has been proposed as an alternative to HDLs as design entry tool to combine high level languages (such as C or C++) with concurrency primitives, thus allowing even faster design cycles for FPGAs than what is now possible using traditional HDLs [119, 189, 39, 229].

Table 3.6 shows some of the commercial software tools currently available in the market.

**Table 3.6.** High Level FPGA Programming Software

Vendor	Product	Base Language
Celoxica	Agility Compiler	Handel-C
Mentor Graphics	Catapult C	C
Impulse Accelerated Tech.	Impulse C	C
Annapolis Microsystems	Core Fire Design Suite	GUI Design Entry
Open System C Initiative (OSCI)	SystemC	C++, IEEE standard 1666

In other order of ideas, designing a complex system in FPGAs can be greatly alleviated by using existing pre-designed libraries. Those libraries, frequently called IP (Intellectual Property) cores, have been fully tested and optimized for performing commonly used building blocks, such as large multiplexers, counters, divisors, digital filters and so forth.

## 3.5 Implementation Aspects for Reconfigurable Hardware Designs

### 3.5.1 Design Flow

In general, most FPGA design tools consist of six basic steps [390] as shown in Fig. 3.9. Those steps must not be executed in a specific order but they can




Fig. 3.9. Design flow

be repeated to improve design's performance. A short description of each step is provided below.

1. **Design Entry** : There are two standard ways to specify an FPGA design, namely,
  - Design Entry through HDLs (Hardware Description Languages): A designer can describe an FPGA design in high-level abstract language like VHDL (Very high speed integrated circuit Hardware Description Language) or Verilog. Those languages are ideal to build state machines, combinational logic, complex and large designs. Most software tools have sophisticated compilers that can efficiently translate HDL specifications to FPGA hardware resources.
  - Design Entry through Schematic: An FPGA design can also be described by using library components of the devices through a graphical interface. It is easy to optimize a circuit for speed/area and consequently it saves time and efforts of the design tool in hardware mapping, placement and routing, etc. However, it is hard to debug and modifications to the design are not straightforward as compared to design entry through HDLs.
2. **Functional verification and simulation:** In this step, the logical correctness of an FPGA design is validated. Once that the design has been specified, either by using HDLs or schematic design entry, it is necessary to verify if such description meets the design specifications.
3. **FPGA synthesis:** Synthesis converts a design entry specification into gates/blocks of an FPGA device. A netlist of basic gates is prepared from HDL/schematic design entry, which is further optimized at gate level. The next step is to map that netlist into FPGA real resources. This is an important step based on design entry. When writing HDL code or using

schematic device's libraries, an FPGA designer should always take into account the basic structure of the target device.

4. **FPGA place and route:** Place and route selects the optimal physical positioning of elementary design blocks and minimal interconnection distance among them. Place and route tools normally use device vendor specifications. Usually they provide hand-placement and also automatic features for optimizing critical paths either for speed or for area.
5. **Circuit analysis:** Circuit analysis evaluates different design performance metrics. Timing verification is made which may differ from functional simulation as it provides logical correctness taking into account all circuit delays occurring in the real device. Similarly, a power analysis evaluation provides an estimation of the design power consumption.
6. **Programming FPGA device:** Programming FPGA implies downloading bit stream codes from the last design steps onto the target FPGA device. Universal programming tools work with FPGAs from different vendors. However there are dedicated programming tools bounded only with a single family of FPGA devices.

### 3.5.2 Design Techniques

It has been observed that better design techniques for both design entry and design implementation play a crucial role for optimizing circuit's performance. A short description of some of those optimizing techniques is given below.

#### Design Strategy

Design strategy is application dependent. For some time critical applications, timing performance is the most important requirement regardless other factors such as hardware resources or device cost. On the contrary, other applications may require a design architecture as compact as possible or with a certain functionality.

Block cipher cryptographic algorithms have an iterative nature, where  $n$  iterations (or rounds) having the same functionality must be executed. It is therefore possible to implement either just one round and consume  $n$  cycles (iterative looping), or  $n$  rounds of the algorithm (using a pipeline structure) in order to achieve high timing performances. The designer choice will be made depending on design's minimum requirements in terms of speed and area.

Fig. 3.10 shows a basic methodology usually followed when implementing an FPGA design.

#### Choice of Target Device

Choosing the target device (FPGA) depends on the design strategy. As it is shown in Table 3.1, an ample spectrum of FPGA devices are available in




Fig. 3.10. Hardware Design Methodology

the market from various manufacturers. The basic structure of all FPGAs is similar, however some models offer additional features like built-in-memories, built-in-arithmetic functions, etc. As it is shown in Table 3.2 for Xilinx devices, different functionality and sizes are available depending on the device's cost.

For example, in the case of block cipher designs it may be useful to select an FPGA device that has embedded Block RAMs (BRAMs) on it. As it was explained above, BRAMs are fast access memories and might be excellent choices for a straightforward implementation of the characteristic S-box blocks of symmetric ciphers. Alternatively, S-Boxes can be implemented using the FPGA CLB fabric configured in memory mode.

In short, the selection of an FPGA depends upon the design size and design requirements.

### Design Analysis

Design/algorithm analysis helps reducing the design's size and critical path delays. It might not be a good idea to directly implement a fast software code in hardware. Software codes are often optimized for high granularity processors, for example, 8, 16 or 32 bit general-purpose microprocessors. Due to its inherent low granularity, hardware implementations quite often can benefit from a bit-level parallelism only limited by data dependencies or resource limitations. For instance, let us consider an instruction from a software code optimized for a 32-bit word-size general-purpose microprocessor:

$$\text{work} = [(((\text{left} \gg 16) | \text{right}) \& 0 \times 0000FFFF)];$$

That requires 16 right shifts, one logical XOR and then one logical AND with  $0 \times 0000FFFF$ . In software platforms, we have no option but to execute an XOR operation for the 16 most significant bits of 32-bit 'left' and 'right' registers.


On the contrary, in hardware description languages, the same instruction can be implemented almost for free, just caring for language notations. One

of the best options is to eliminate the AND operation and 16 logical Shifts by executing instead an XOR operation directly applied to the 16 most significant bits of left and right registers, that is,

$$\text{work} = \text{left}[31:16] \oplus \text{right}[31:16]$$

### Selecting FPGA Resources

An FPGA designer can pick multiple options for performing a function. For example, two choices for implementing a 2-bit multiplexer are shown in Figure 3.11.



**Fig. 3.11.** 2-bit Multiplexer Using (a) Tristate Buffer. (b) LUT

Figure 3.11.a shows usage of tri-state buffers for a multiplexer. A large number of tri-state buffers are available in FPGAs and it seems logical to make use of them. However, experience shows that, using large number of tri-state buffers slows down the circuit. This tends to require the physical distribution of tri-state buffers all around FPGA, which requires long routing paths. A multiplexer can also be implemented using LUTs as shown in Figure 3.11.b. Using adjacent LUTs for an  $n$  to 1 multiplexer would be useful when a circuit must be optimized for speed.

Similarly, some FPGA devices contain built-in memory modules. It would be useful to utilize those memories as they provide faster access to the data as compared to distributed memories in FPGAs which are formed using several LUTs.

### Hardware Approach

A careful selection and usage of the design tools results useful in our methodology for obtaining better performances. The design tools by Xilinx [390], Altera [3], Synopsis Galaxy Design Platform [351], LeonardoSpectrum and ModelSim by Mentor Graphics [231, 230], etc. provide several useful features for getting design improvements. Better placement of the components or better routing of the architecture modules can be helpful in cutting critical path delays in the circuit.

3.5.3 Strategies for Exploiting FPGA Parallelism

Achieving high-speed implementations for cryptographic algorithms is an exciting task requiring deep considerations at every stage of the design. Design strategies should therefore not only be based on the best implementing techniques on reconfigurable platforms but also on trying to innovate in the theoretical side by improving the standard transformations of cryptographic algorithms. In this sense, the designs included in this book try to take as much advantage as possible of the hardware inherent parallelism while keeping as low as possible the hardware resource requirements. In the following we discuss various strategies used by designers to implement cryptographic algorithms.

Iterative Looping (IL)

An iterative looping design (IL), implements only one round and  $n$  iterations of the algorithm are carried out by feeding back previous round results as shown in Figure 3.12a. It utilizes less area but consumes more clock cycles resulting on a relatively low speed encryption.

Loop Unrolling

Architecture with loop unrolling is shown in Figure 3.12b. In a loop unrolling or pipeline design (PP), rounds are replicated and registers are provided between the rounds to control the flow of data. The design offers high speed but area requirements tend to be too high.

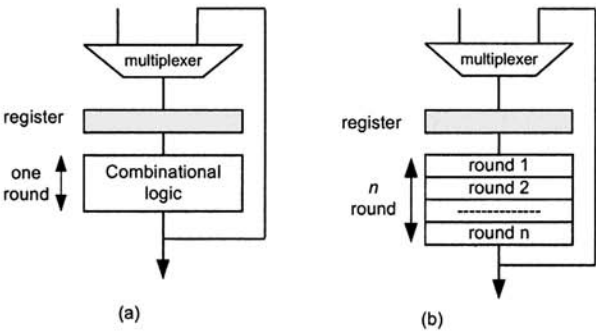


Fig. 3.12. Basic Architectures for (a) Iterative Looping (b) Loop Unrolling



### Inner-Round Pipelining

Figure 3.13a shows an inner-round pipelining architecture where extra registers are provided at different stages of the same round in such a way that several blocks of data can be processed by the circuit at the same time. This approach produces high speed circuits at the cost of more hardware resources in the form of registers.

### Outer-Round Pipelining

Outer-round pipelining is created through loop unrolling by adding extra registers at different stages of the same round as shown in Figure 3.13b. This approach directly trades circuit speed with circuit area.




Fig. 3.13. Round-pipelining for (a) One Round (b)  $n$  Rounds

Both the iterative and pipeline architectures would be optimized for the implementation of secret-key ciphers. Public key algorithms exhibit different nature. They do not have rounds however they maintain a hierarchical structure that can be further exploited.

## 3.6 FPGA Architecture Statistics

Just as it occurs with software platform comparisons, comparing FPGA designs is a difficult and a bit ambiguous task. The two single most important performance metrics usually considered are the *time complexity*, sometime called design throughput and the *area complexity*.

For combinatorial designs (such as adders, squarers, fully-parallel multipliers, etc), time complexity is determined from the Maximum Clock Frequency

(MCF), which in turns is proportional to the maximum combinational path delay. In the case of sequential designs (such as block ciphers, sequential multipliers, etc.), time complexity must also consider the total number of clock cycles required before the result is ready. In the case of block cipher designs, it is customary to consider also how many bits are processed at the same time. In this work we define the throughput of a given design as follows,

### Throughput

Throughput is an important factor to measure timing performances of the design [82, 103, 382]. Throughput of the design is obtained by multiplying the allowed frequency for the design with the number of bits processed per cycle. For cryptographic algorithms, throughput is defined as:

$$\text{Throughput} = \frac{\text{Allowed Frequency} \times \text{Number of Bits}}{\text{Number of cycles}} \text{ (bits/s)}$$

The higher the throughput of a design is the better its efficiency.

### Area

Design statistics provided by the design software expresses hardware area occupied by the design. Unfortunately, there is no universal metric to measure the hardware costs associated with an FPGA based design. After mapping a design to a particular FPGA device, FPGA compiler provides FPGA resources utilized by that design.

Following are some common FPGA resources listed by the mapping tool:

- Number of slices
- Number of Slice Flip Flops
- Number of 4-input Look Up Tables (LUTs)
- Number of Input/Output Blocks
- Number of Clocks
- Maximum combinational path delay
- Maximum output required time after clock
- Maximum Clock Frequency (MCF)
- BlockSelect RAMs (BRAMs)

A designer, however, can report hardware area in terms of LUTs as well as CLB slices. An ideal comparison would be therefore comparing all resources on the similar FPGA device. A design using dedicated resources of the device will show less logic resources as compared to other design which implements the whole logic without using any dedicated unit of the device. It also affects the throughput statistic. It has been experimentally observed that the implementation of even the same code on different grades of the same family of devices influence the final design's throughput. That situation becomes more

crucial when the same design targets two different devices by two different manufactures. In such cases, for the purpose of classifying an FPGA design, we can ignore some of those factors.

It can be said, as a first approximation, that the fastest design is the one which achieves fastest speed no matter what type of device has been targeted for design implementation. However, when considering a compact design (a design optimized for hardware area), this criterion cannot be applied. The comparison of two compact designs can be only justified if it is made between similar devices.

Both area and throughput factors provide a measure for comparing different designs. Additionally, in order to decide how efficient a design is, we utilize the following figure of merit.

### **Throughput/Area**

It is the ratio of the above two figures of merits and shows how efficient the design is with respect to both area and throughput. The ratio is higher in case of high throughput and less space.

## **3.7 Security in Reconfigurable Hardware Devices**

The selection of an implementation platform in a digital system depends on many design criteria. Besides the design performance figures such as, system speed and area costs, there exist other performance and security factors that should be taken into account such as: physical security (for instance, against key recovery and algorithm manipulation), flexibility, power consumption and other secondary factors, that may as well affect the design selections.

Even though there exist a fair amount of papers reporting cryptographic implementations on FPGA devices, there are not that many papers reporting the convenience (or not) of utilizing FPGA as a target device for security applications from a system point of view. In particular, few works report the resilience of FPGA against physical or system attacks, which are potentially more dangerous than algorithm attacks [379, 342, 343].

In [380, 379] a comprehensive analysis of FPGA security aspects is given. Authors conclude that FPGA technology can provide a reasonable level of security when used properly.

The fourth generation design security of Xilinx Virtex-4 family is equipped with bit-stream encryption/decryption technology based on 256-bit AES. The user generates the encryption key and encrypted bit-stream using Xilinx ISE software. In a second step, during configuration, the Virtex-4 device decrypts the incoming bit-stream using a decryption logic module with dedicated memory for storing the 256-bit encryption key [393].

For the cryptographic applications, the most important threat is unauthorized access to a confidential cryptographic key, either a symmetric key or the private key of an asymmetric algorithm<sup>4</sup>.

FPGA implementations are also vulnerable to side-channel attacks. A side channel attack is based on information gained directly from the physical implementation. Examples for side channels include: power consumption, timing behavior, and electromagnetic radiation. Most relevant papers on side-channel attacks and related defenses have been published in [183, 184, 182, 159, 366, 157, 278].

Power analysis attacks were introduced in 1998 by Kocher et al. [186]. The main idea behind this attack is to measure the power consumption of the FPGA device during the execution of a cryptographic operation. Thereafter, that power consumption can be analyzed in an effort for finding regions in the power consumption trace of a device that are correlated with algorithm's secret key.

In [262], the first experimental results of power analysis attack on an FPGA implementation of elliptic curve cryptosystem were presented. RSA, AES and DES FPGA implementations have also been subjects of attacks in [341, 342, 343].

### 3.8 Conclusions

In this chapter we presented some of the most relevant aspects related to FPGA devices considering both, technological and reconfigurable programming aspects.

The material covered in this Chapter includes a brief review of the technological antecedents that gave birth to FPGA devices. We also studied the structure of several emblematic FPGA families from the two market leaders, Xilinx and Altera. We compare the performance of FPGA realizations against the ones on ASICs and general-purpose processor platforms and we briefly introduced the main concepts related to the reconfigurable computing paradigm.

Furthermore, we reviewed several key strategies to achieve good designs when working with cryptographic applications. As a way to measure area and time performances for a given design, we defined several metrics and figures of merit. Finally, several security concerns related to FPGA technology were outlined.

---

<sup>4</sup> As it was described in the precedent chapter, most cryptographic algorithms have been standardized and therefore, they are publicly known.